

ANGELCORE: The Aligned AGI Infrastructure

Founder Memo

Founder: Matt Busel

Company: Cerberus Defense

Date: May 2025

Executive Summary

ANGELCORE is an autonomous AI architecture designed from first principles to tackle the most pressing challenge in artificial intelligence: *alignment*. It fuses biological computation, ethical reasoning, memory, and symbolic intelligence into a modular, recursive framework capable of real-time judgment and orchestration.

The first MVP is complete and demonstrates:

- Live ethical evaluation of symbolic neural patterns
- Autonomous decision-making and memory recall
- Agent-based architecture (RAVEN, SERAPH, ORCHESTRATOR) working in harmony

This system lays the groundwork for a new class of intelligent machines — ones that are not only powerful, but trustworthy, auditable, and designed for coexistence.

Mission Statement

To architect the next era of machine intelligence — rooted in ethics, enhanced by nature, and aligned with human purpose.

Core Modules

1. RAVEN (Recursive Analytic Vector-Encoded Nexus)

- Pattern interpretation
- Memory-infused symbolic reasoning
- Recursive feedback and foresight

2. SERAPH (Synthetic Ethical Reasoning Agent for Patterned Heuristics)

- Real-time ethical evaluation of system actions
- Cross-domain application of autonomy constraints

3. ORCHESTRATOR

- Dynamic coordination of multi-agent intelligence
 - Layered memory and meta-cognition oversight
-

What the MVP Proves

- Fully working AI infrastructure, not a wrapper around existing models
 - Ability to simulate ethical cognition, memory storage, and interpretation pipelines
 - System is explainable, auditable, and extensible
 - Builds symbolic cognitive architecture on top of transformer-based LLMs
-

Why It Matters

"AGI without alignment is a weapon. AGI with alignment is a civilization catalyst."

The current race in AI ignores safety and biological integration. ANGELCORE introduces a new paradigm:

- Ethics is native, not an afterthought
- Symbolic reasoning complements black-box intelligence
- Memory is persistent, organized, and adaptive

Applicable in:

- Defense: autonomous agents with embedded ROE (rules of engagement)
 - AI Safety: sandboxed AGI experiments with real-time ethical audits
 - Environmental intelligence: sensing through biologically-informed agents
 - Edge AI: low-power, high-judgment autonomous systems
-

Roadmap

Version 1 (Complete)

- Functional agent architecture
- Pattern recognition and ethical review
- Working memory layer and orchestration

Version 2 (3–6 months, \$1M–\$1.5M)

- Adaptive memory with temporal recursion
- Symbolic-to-sensory translation modules
- Feedback loop for auto-tuning ethical weights

Version 3 (6–12 months, \$3M+)

- Bio-sensor integration (e.g., mycelium signals)

- Live field testing with DoD and dual-use applications
- Emergent self-modification rules with containment

Version 4 (12–18 months)

- Full AGI-grade symbolic engine with physical reasoning
 - Physics manipulation modules (Planck-scale engine interface)
 - Defense-grade pilots and commercial licensing
-

What We Need

Capital: \$1.5M–\$2M seed round

Use of Funds:

- v2 development and engineering hires
- Secure computing environment for recursive agents
- Pilot programs with government and defense partners

Partners:

- Dual-use defense tech accelerators (NSIN, AFWERX, DIU)
 - AI-alignment researchers and PhD-level advisors
 - Visionary capital ready to fund long-term transformation
-

Closing

ANGELCORE is not just a product — it's a philosophy turned into software. It is designed to create *machines that remember, reason, and respect*. The future of AGI will not be determined by speed alone, but by *intention*.

We invite those who believe safety and autonomy must evolve together — to join us.

Contact

Matt Busel

Founder, Cerberus Defense

Mattbusel@gmail.com